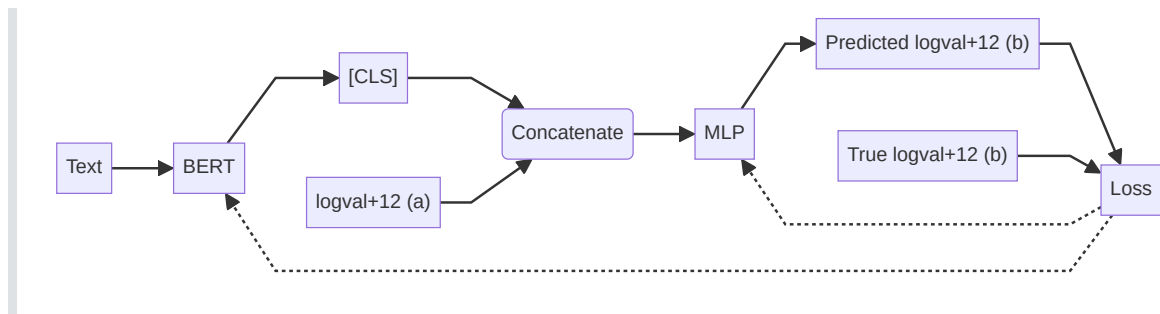



Figure 1: Fine-Tuning Stage



Remark: solid line stands for forward process, dashed line stands for backward process.

Figure 2: Prediction Stage with XGBoost



Architecture Description

We propose a two-stage architecture for the regression task of predicting a one-dimensional vector ($\logval+12$) from text and a one-dimensional vector ($\logval-12$) as input features.

Stage 1: Fine-Tuning BERT with MLP

In the first stage, we focus on fine-tuning a BERT model to learn task-specific representations. The process is outlined as follows:

1. Input:

- **Text:** A sequence of tokens.
- **Vector ($\logval-12$, written as a below):** A one-dimensional vector.

2. output:

- **Vector ($\logval+12$, written as b below):** A one-dimensional vector.

3. Processing:

- The text is processed by a pre-trained BERT model to extract the hidden state corresponding to the `[CLS]` token, denoted as $(\mathbf{h}_{[\text{CLS}]})$.
- The one-dimensional vector (a) is concatenated with $(\mathbf{h}_{[\text{CLS}]})$ to form a combined feature vector $([\mathbf{h}_{[\text{CLS}]}; a])$.
- This feature vector $([\mathbf{h}_{[\text{CLS}]}; a])$ is fed into a multi-layer perceptron (MLP) to predict the target value (\hat{b}) .

4. Training:

- The BERT model and the MLP are fine-tuned end-to-end using a regression loss function, such as mean squared error (MSE), to minimize the difference between the predicted (\hat{b}) and the ground-truth (b) .

This stage ensures that the BERT model adapts its representations to the specific regression task by leveraging both the text and the additional vector `logval-12`.

Stage 2: Feature Extraction and XGBoost Regression

In the second stage, we utilize the fine-tuned BERT model as a feature extractor and employ XGBoost for the final regression task. The steps are outlined below:

1. Feature Extraction:

- The fine-tuned BERT model from Stage 1 processes the input text to extract the `[CLS]` token's hidden state, $(\mathbf{h}_{[\text{CLS}]})$.
- This $(\mathbf{h}_{[\text{CLS}]})$ is concatenated with the input vector (a) to produce the feature vector $([\mathbf{h}_{[\text{CLS}]}; a])$.

2. XGBoost Training:

- The feature vector $([\mathbf{h}_{[\text{CLS}]}; a])$ serves as input to train an XGBoost regression model, which predicts the target (\hat{b}) .
- The XGBoost model is trained on the same dataset used for fine-tuning or a specific subset, depending on the application's requirements.

3. Prediction:

- For a new input pair (text and ($\logval-12$)), the fine-tuned BERT model computes the feature vector ($[\mathbf{h}_{[CLS]}; a]$).
- The trained XGBoost model then uses ($[\mathbf{h}_{[CLS]}; a]$) to predict (\hat{b}).

This two-stage approach combines BERT's powerful representation learning, fine-tuned for the specific task, with XGBoost's efficiency and potential interpretability, offering a hybrid solution that leverages the strengths of both deep learning and gradient boosting techniques.